


# Lab 04 - Becoming a Scheduler



**Instructor:** Lorenzo De Carli, University of Calgary ([lorenzo.decarli@ucalgary.ca](mailto:lorenzo.decarli@ucalgary.ca))

*Slides by Lorenzo De Carli, based on material by Robert Walls (WPI)*

# In this lab

- Build a **scheduler simulator**
  - Receive as input a **sequence of jobs** in a file
  - **Simulate execution** of those jobs
  - Output an **execution trace** with:
    - Job start time
    - Job end time

# Deliverable

## A scheduler executable

- You will need to provide a **Makefile** to compile your code
- Compiler should produce a **scheduler.out** file accepting **3 parameters**:
  - A flag (**0** or **1**) detailing whether or not to perform **policy analysis**
    - ...more details about this later
  - Name of the scheduling policy (**FIFO** or **SJF**)
  - Name of input job file (e.g. **jobs.txt**)

# Example

## Running the SJF policy w/ no analysis from jobs.txt

`./scheduler.out 0 SJF jobs.txt`

- You can assume **all parameters** are **always specified**
- Here **0** means “no policy analysis”
- **SJF** means to run the SJF scheduling policy
- **jobs.txt** means to read the list of jobs from the file “jobs.txt”

# Input file format

- Each **workload** is defined in a **workload file**.
- Each line of the workload file represents a **different job** in the workload
- Each line consists of **two comma-separated numbers**:
  - the **arrival time**, and
  - the **total amount of simulated time** that job needs to run.

# Input file format - example

0,1

2,5

3,11

5,4

7,1

25,4

← Each line includes **arrival time** and **duration**

You can assume there are **no extra spaces**,  
and that **every line ends with a newline**

# So... what's about these jobs?

- To be clear, those are **not actual jobs**
- The scheduler should **simulate** that sequence of jobs by appropriately **computing when they start and end**
- There is no need to actually run anything (except the **scheduler calculations**)

# Job list data structure

- The scheduler uses the job file to initialize a **job list data structure**
- In practice, this should be a **linked list**
- Each job should be assigned an **id** based on the **line number in the file**
- The job on the **first line** should be assigned an id of **0**; the job on the **second line** should be assigned an id of **1**; and so on



# Job list data structure /2

This is just an example...

```
struct job {  
    int id;  
    int arrival;  
    int length;  
    // other meta data  
    struct job *next;  
};
```

# Implementing policies

# Exercise 1: Implementing FIFO

- The **FIFO policy** is one of the simplest scheduling policies
  - Good starting point! 😊
- The FIFO policy states that **jobs are scheduled in order of their arrival**
- Each job **runs to completion**
- To be clear: there is **no preemption** for this FIFO policy.

# Example scheduler output...

## ...when running FIFO

```
$ ./scheduler.out 0 FIFO tests/3.in
```

Execution trace with FIFO:

t=0: [Job 0] arrived at [0], ran for: [20]

t=20: [Job 1] arrived at [0], ran for: [19]

t=39: [Job 2] arrived at [1], ran for: [18]

t=57: [Job 3] arrived at [1], ran for: [17]

t=74: [Job 4] arrived at [2], ran for: [16]

t=90: [Job 5] arrived at [3], ran for: [15]

t=105: [Job 6] arrived at [4], ran for: [14]

End of execution with FIFO.

# Exercise 2: Implementing SJF

(aka “Shortest Job First”)

- **SJF** always picks the job with the **shortest runtime** to run next
- We again assume that a job will run to completion **before the next is started**
- If two jobs need the same amount of time, SJF breaks the tie by favoring **the job that arrived earlier**
- Your SJF scheduler should account for periods when there are **no jobs** to run
  - That is, all the arrived jobs **have completed** before the new jobs arrive
  - In other words, **the CPU can be idle**

# Example scheduler output...

## ...when running SJF

\$ ./scheduler 0 SJF tests/[8.in](#)

Execution trace with SJF:

t=0: [Job 0] arrived at [0], ran for: [1]

t=2: [Job 1] arrived at [2], ran for: [5]

t=7: [Job 4] arrived at [7], ran for: [1]

t=8: [Job 3] arrived at [5], ran for: [4]

t=12: [Job 2] arrived at [3], ran for: [11]

t=25: [Job 5] arrived at [25], ran for: [4]

End of execution with SJF.

Note that the CPU was  
**idle** for 1 tick here



# Policy analysis

# Policy analysis?

- In this part of this project, you will add code to the scheduler to help it evaluate the **performance** of the previously implemented **policies**
- Your code will measure **two metrics**:
  - Response time
  - Turnaround time



# Metric definitions

- Assume:
  - $T_a$  is the job **arrival time**
  - $T_s$  is the job **start time**
  - $T_c$  is the job **completion time**
- Then:
  - **Response time** is  $T_s - T_a$
  - **Turnaround time** is  $T_c - T_a$

# Exercise 3: Scheduler policy analysis

- The modified scheduler should output, for each metric:
  - The **per-job value** of the metric
  - The **average value** of the metric **across all jobs**

# Example FIFO scheduler output...

## ...with metrics

```
$ ./scheduler 1 FIFO tests/3.in
```

Execution trace with FIFO:

t=0: [Job 0] arrived at [0], ran for: [20]

t=20: [Job 1] arrived at [0], ran for: [19]

t=39: [Job 2] arrived at [1], ran for: [18]

t=57: [Job 3] arrived at [1], ran for: [17]

t=74: [Job 4] arrived at [2], ran for: [16]

t=90: [Job 5] arrived at [3], ran for: [15]

t=105: [Job 6] arrived at [4], ran for: [14]

End of execution with FIFO.

Begin analyzing FIFO:

Job 0 -- Response time: 0 Turnaround: 20 Wait: 0

Job 1 -- Response time: 20 Turnaround: 39 Wait: 20

Job 2 -- Response time: 38 Turnaround: 56 Wait: 38

Job 3 -- Response time: 56 Turnaround: 73 Wait: 56

Job 4 -- Response time: 72 Turnaround: 88 Wait: 72

Job 5 -- Response time: 87 Turnaround: 102 Wait: 87

Job 6 -- Response time: 101 Turnaround: 115 Wait: 101

Average -- Response: 53.43 Turnaround 70.43 Wait 53.43

End analyzing FIFO.

# Example SJF scheduler output...

## ...with metrics

```
$ ./scheduler 1 SJF tests/8.in 0 Execution trace with SJF:
```

```
t=0: [Job 0] arrived at [0], ran for: [1]
```

```
t=2: [Job 1] arrived at [2], ran for: [5]
```

```
t=7: [Job 4] arrived at [7], ran for: [1]
```

```
t=8: [Job 3] arrived at [5], ran for: [4]
```

```
t=12: [Job 2] arrived at [3], ran for: [11]
```

```
t=25: [Job 5] arrived at [25], ran for: [4]
```

```
End of execution with SJF.
```

```
Begin analyzing SJF:
```

```
Job 0 -- Response time: 0 Turnaround: 1 Wait: 0
```

```
Job 1 -- Response time: 0 Turnaround: 5 Wait: 0
```

```
Job 2 -- Response time: 9 Turnaround: 20 Wait: 9
```

```
Job 3 -- Response time: 3 Turnaround: 7 Wait: 3
```

```
Job 4 -- Response time: 0 Turnaround: 1 Wait: 0
```

```
Job 5 -- Response time: 0 Turnaround: 4 Wait: 0
```

```
Average -- Response: 2.00 Turnaround 6.33 Wait 2.00
```

```
End analyzing SJF.
```

# Code template

## File-by-file description

- **scheduler.c**: template file to complete to implement the scheduler
- **example\_fifo.in**: example input to test the FIFO scheduler
- **example\_fifo.out**: expected output when running `./scheduler.out 0 FIFO example_fifo.in`
- **example\_fifo\_analysis.out**: expected output when running `./scheduler.out 1 FIFO example_fifo.in`
- **example\_sjf.in**: example input to test the SJF scheduler
- **example\_sjf.out**: expected output when running `./scheduler.out 0 SJF example_sjf.in`
- **example\_sjf\_analysis.out**: expected output when running `./scheduler.out 1 SJF example_sjf.in`

# More on the code template

- The `readme.md` file (also in the template directory) contains useful information on **how to test your code**
- Most importantly, **don't assume we are going to use the reference inputs/outputs when grading.**
- If your code works on the sample inputs/outputs but **fails on other ones**, you will still **lose points**
- See `readme.md` for some hints on how to **test your work efficiently**
- Finally, **remember to create a Makefile**

# Grading rubric

...the part everyone cares about!

- As usual, you get **3 pts** for uploading a **partial solution** by **end of lab (1 PM)**
- Then, you'll have until **11:59PM of the day before the next lab** to upload the **complete solution**. That will be graded as follows:
  - Correct **FIFO implementation**: 2 pts
  - Correct **SJF implementation**: 2 pts
  - Correct **FIFO analysis**: 1.5 pts
  - Correct **SJF analysis**: 1.5 pts