

Lab 06 - Memory Allocation

Instructor: Lorenzo De Carli, University of Calgary (lorenzo.decarli@ucalgary.ca)
Slides by Lorenzo De Carli, based on material by Robert Walls (WPI)

Introduction

- Have you ever wondered what happens when you call **malloc()**? In this lab you will figure it out (sorta)
- You will have **two weeks** to complete this lab. As such, you may find it amore challenging than the previous one
- You are encouraged to **read chapter 17 of the book**. We won't cover this material in class, but it will provide a lot of **useful context** for the lab

What is the goal?

...to implement a memory allocator!

- You must implement a memory allocation library in two files (**myalloc.h** and **myalloc.c**)
- The library implements four functions:
 - **void myinit(size_t size);** // initialize the memory allocator
 - **void* myalloc(size_t size);** // request allocation of new memory
 - **void myfree(void* ptr);** // free memory allocated with 461alloc
 - **void mydestroy();** // release all resources used by the allocator

**Warning: you are entering a
malloc()-free zone**

**At no point in this assignment
you are allowed to use malloc()**

Function prototypes

- The four functions are prototyped in the header file provided in the code template. The header is called **myalloc.h**
- **Do not modify this file**, otherwise you will get **0 points** for the lab
- Your code should go into **myalloc.c**
- To avoid **compiler errors**, please make sure your implementation of **myalloc.c** includes the following:

```
#include <stddef.h>
```

```
#include "myalloc.h" // This include should appear last
```

Test cases

- The assignment is divided in **5 conceptual parts**
- We will provide **test cases for each part** in the code template
- You can run part-specific test cases by going into the appropriate folder and run **test-partX.sh**
- There is a **test_all.sh** script you can use to run the test cases for all parts
- For obvious reasons, we suggest that you make sure that your code **passes all test cases**

Part 1

Initialization and Destruction

- **myinit()** must get a large memory area (called an **arena** in the following) from the OS, which will be used to allocate the chunks of memory requested by **myalloc()**
- How do I request such a contiguous area? The easiest is to use **mmap()**. Your code should look like the lines below:

```
_arena_start = mmap(NULL, size, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0);
```

Example `myinit()` output

Initializing arena:

- ...requested size 1 bytes
- ...pagesize is 4096 bytes
- ...adjusting size with page boundaries
- ...adjusted size is 4096 bytes
- ...mapping arena with `mmap()`
- ...arena starts at `0x7f8e5f4f9000`
- ...arena ends at `0x7f8e5f4fa000`

What is this `mmap()` business?

...just the most confusing syscall ever

- If you are confused about `mmap()`, then you are not alone
- The canonical use of `mmap()` (and probably its original purpose) is to **map a file to a memory region**
 - Reads/writes to **memory** are automatically transferred to the underlying **file**
- However, it also allows to **allocate a memory region without mapping it to any file** (that's how you are using it here)
- There are reasons to use it which we won't get into - but it can be **more efficient** than `malloc()` for large allocations

What about destruction?

- **mydestroy()** should de-allocate resources allocated by **myinit()**
- Primarily, it should use **munmap()** to de-allocate the memory requested by **mmap()**
- **mydestroy()** should also reset any state variable that your code keep, restoring them to their original value
- “**man mmap**” and “**man munmap**” are your friends here! 😊

Example mydestroy() output

Destroying Arena:

...unmapping arena with munmap()

Part 2

...the actual allocation

- The **myalloc()** function can be called by program code to request a memory allocation (exactly like **malloc()**)
- The allocated memory must reside inside the allocated arena
- Once allocated, the program should be able to read/write to that memory until it is freed using **myfree()**
- When **myalloc()** is called, your code will allocate a memory chunk. Let's see what is inside...

Allocation chunk

- A **chunk** of allocated memory should consist of **two parts**:
 - A **header** containing metadata about the chunk, of type **node_t** (already defined in **myalloc.h**)
 - The actual area of memory **allocated to the application**
- Notes:
 - The size of the chunk should account for the **space taken by the header**
 - Make sure you **do not return the header** to the program!

What about freeing?

- **myfree(void* ptr)** should free the memory allocated by myalloc()
- Note that **ptr** will not point to the header - it will point to the part of a chunk which is **reserved to the application!**
- If the program passes an address **a** to **myfree()**, to get the address of the chunk header, you will need to do something like:

((void *)a) - sizeof(node t)

Putting it all together

- Chunk headers of type **node_t** must be organized in a **doubly-linked list**
- In practice this means:
 - The chunk headers are allocated **within each chunk**
 - Each chunk header has pointers to the **next** and the **previous chunk**
- You must keep chunk headers both for **allocated** and **free** chunks of memory

Warning: you are in a malloc()-free zone

**At no point in this assignment
you are allowed to use malloc()**

Some graphical examples


- Here is how an arena of size 100 looks after initialization

```
Chunk 0  
size: 100  
is_free: 1  
fwd: NULL  
bwd: NULL
```

Chunk header

Chunk memory

- Here is what happens after **myalloc(20); myalloc(30); myalloc(10)**



- Here is what happens after **myfree(chunk1)**

Chunk0 size: 20 is_free: 0 fwd: Chunk1 bwd: NULL	Chunk1 size: 30 is_free: 1 fwd: Chunk2 bwd: Chunk0	Chunk2 size: 10 is_free: 0 fwd: Chunk3 bwd: Chunk1	Chunk3 size: 40 is_free: 1 fwd: NULL bwd: Chunk3
--	--	--	--

Note: if the program requests more memory than available, **myalloc()** should return null and set **statusno** to **ERR_OUT_OF_MEMORY**

Example output of myalloc():

Allocating memory:

- ...looking for free chunk of ≥ 4064 bytes
- ...found free chunk of 4064 bytes with header at 0x7f9655b92000
- ...free chunk→fwd currently points to (nil)
- ...free chunk→bwd currently points to (nil)
- ...checking if splitting is required
- ...splitting not required
- ...updating chunk header at 0x7f9655b92000
- ...being careful with my pointer arithmetic and void pointer casting
- ...allocation starts at 0x7f9655b92020

Example output of myfree():

Freeing allocated memory:

- ...supplied pointer 0x7fd515872120:
- ...being careful with my pointer arithmetic and void pointer casting
- ...accessing chunk header at 0x7fd515872100
- ...chunk of size 3808
- ...checking if coalescing is needed
- ...coalescing not needed.

Part 3

Free chunk splitting

- Suppose a requested allocation is **smaller than free chunk size** (e.g., `myalloc(10)` but the free chunk has size 25)
- In that case, the allocator must **split the chunk in two**: one allocated chunk of the requested size, and one free chunk of the remaining size
- I slightly lied to you above, because the actual size of the chunk is *requested space + size of the chunk header*
 - **Make sure you take that into account!**

Part 4

Placement of allocations

- The chunk list must be **logically ordered by address** of each chunk, in **increasing order**
- The **first chunk** in the list will be the chunk with the **lowest address**, and the **last chunk** will be the one with the **highest address**
- The ordering should be **naturally maintained** by properly implementing **myalloc()** and **free()**; in other words you should not need to implement any sorting function
- Your allocator should satisfy a **myalloc()** request by **allocating the first free chunk** which is **large enough** for the request

Part 4 example

- The following sequence should result in **buff3** being placed at the previous location of **buff1**:

```
buff1 = myalloc(64);  
buff2 = myalloc(64);  
myfree(buff1);  
buff3 = myalloc(48);
```

Part 5

Free chunk coalescing

- Suppose freeing a chunk results in two or more **contiguous free chunks** in the list
- Then, your allocator should **coalesce those chunks** (i.e., merge them) into a **single larger chunk**

Part 5 example

- The following sequence should result in the memory allocator coalescing all three chunks in a single free chunk when the last **myfree()** is called:

```
buff1 = myalloc(64);  
buff2 = myalloc(64);  
buff3 = myalloc(64);  
myfree(buff1);  
myfree(buff3);  
myfree(buff2);
```

Test cases

...find them in the code template!

- We are going to grade your code based on these test cases
- Hardcoding output to pass test cases will result in a **score of 0** and the student(s) being **referred for academic misconduct**

How to run tests?

- First, you may need to give execution permission to all test scripts:

```
chmod a+x test_all.sh
cd test-part1
chmod a+x test_part1.sh
cd..
cd test-part2
chmod a+x test_part2.sh
cd..
```

- ...and so on

How to run tests? /2

- After that, you can run individual tests by doing (e.g., for part 1):

cd test-part1
./test_part1.sh
- From the output, it will be clear which tests pass and which ones do not
- Look into the **tests.c** files in each test directory to understand what specifically is being tested
- To run all tests, run **./test_all.sh**
- Note, **myalloc.c** should be in the same directory as **myalloc.h** (no need to copy it into **test-part1** or any other subdirectory)

Grading rubric

- You get **3 points** for submitting your initial work during the lab
- The remaining 7 points are distributed as follow (not passing all test cases for one part will result in a score of 0 for that part; we will not assign partial scores):
 - Passing all part 1 test cases: **1 pt**
 - Passing all part 2 test cases: **2 pts**
 - Passing all part 3 test cases: **1 pt**
 - Passing all part 4 test cases: **1 pt**
 - Passing all part 5 test cases: **2 pts**

Deadline & delivery

- Given the upcoming midterm, you get two weeks to complete this lab
- Submit your **final version** on **the day before the next lab**, which is in **two weeks**
- Note, there is **no lab next week**
- You just need to deliver **myalloc.c**. Do not upload **myalloc.h**, test cases, or any other file

Advice

- Reach **Chapter 17 of the book** to get some **context**
- Note, the allocator presented here **works slightly differently** from the one on the book
- Familiarize yourself with **void pointers (`void*`)**, **structs**, and **pointer arithmetics**
 - The **test cases** have **useful examples**
- Familiarize yourself with **`size_t`** and remember that it is **unsigned**
- Always check for **NULL** pointers!

That is all... good luck!