

Lecture 04 - C Refresher

ENSF461 - Applied Operating Systems

Instructor: Lorenzo De Carli, University of Calgary (lorenzo.decarli@ucalgary.ca)

Slides by Lorenzo De Carli

Review of previous lecture

Command-line utilities

- Command-line utilities assist in **automating various tasks**:
 - **Text processing**: filtering and aggregating text data; computing statistics
 - **Operating on the file system**: changing working directory, creating and deleting files and folders
 - Show various machine statistics (e.g., free disk space)

More command line

The `grep` utility

- **`grep`** parses a file (or the standard input) and returns **all lines that match (or do not match) a string**
- A few options:
 - **`-i`**: makes matching case insensitive
 - **`-v`**: return all lines that do not match the pattern
 - **`-c`**: print number of matching lines per file
 - **`-l`**: print names of files with matches

Running commands on multiple files

Let's talk about wildcards

- It is typically to be possible to run a utility on **more than one file** in one pass
- For example, use **grep** to find a string in **all files in a directory**
- This is accomplished using **wildcard patterns**:
 - **\$ grep hello *.c**: find the string “hello” in all files ending in “**.c**” in the current directory
 - **\$ grep hello /home/lorenzo/prog*c**: find the string “hello” in all files starting in “**prog**” and ending in “**c**” in the directory **/home/lorenzo**

Shell programming in Bash

```
#!/bin/sh

VAR=1
while [ $VAR -lt 10 ]
do
    MOD=`expr $VAR % 2`
    if [ $MOD -eq 0 ]
    then
        echo "$VAR is even"
    else
        echo "$VAR is odd"
    fi
    VAR=`expr $VAR + 1`
done
```

The plan for today

Let's meet an old friend

- Why are we going to use **C** for this course?
- **C basic syntax + hello world**
- Simple **data structures**
- Compiling **C programs** from the **command line**
- **Debugging** C programs

A disclaimer

- Learning C is not the topic of this course
- We will provide you a quick refresher in this lecture
- We will also explain how to build and debug programs on Linux VMs
- The rest is up to you



Why C?

Easier to say why not!

- C is very **old!**
 - **Even older than your instructor** 😱
- C is **very unfriendly**
 - It's a lot more work than Python etc.
- C is **dangerous**
 - Easy to make **memory errors...**
 - ...which make it **insecure**

A new report examining the security of programming languages has found that almost 50% of all the vulnerabilities discovered in open source projects since 2009 were coded in C.

The study by WhiteSource revealed that 46.9% of all reported open source vulnerabilities in the past 10 years were developed using C.

<https://portswigger.net/daily-swig/c-is-least-secure-programming-language-study-claims>

So... WHY C?!

There are some good reasons

- C is as close as it gets to the **operating system interface**
 - Except if we programming in **assembly**, which is a bigger **can of worms**
- C was develop in tandem with **UNIX** operating systems
 - Many UNIX (and Linux) interfaces are **straightforwardly accessible** from C
- C is the language used in the **book**
 - Do you really want to spend all of the course porting the examples from the book to **Rust**?

So yeah, we are using C

A note about C++

- **C++** shared many of the properties discussed before...
- ...but is considerably **more complicated** than C
- Also, it hides some of the **low-level details** we care about
- **In class we are mostly going to stick to C**
 - May use some C++ constructs if makes things easier

What is C?

- **Programming language** first proposed by Dennis Ritchie in the early '70s
- Why “C”? It was derived from an earlier language called “B”
- Characteristics:
 - **General-purpose**
 - **Compiled (what does it mean?)**
 - **Strongly-typed (what does it mean?)**

C's basic syntax

There are four program sections you need to worry about

```
#include <stdio.h>
```

} — Preprocessor directives

```
void do_stuff();
```

} — Functions/globals declarations

```
int main() {  
    do_stuff();  
}
```

} — Main function (entry point)

```
void do_stuff() {  
    printf("Hello, world!");  
}
```

} — Function definitions

Let's look at each of them...

Preprocessor directives

What is a “preprocessor”?

- The **C preprocessor** expands **macros** (marked with “**#**”) before sending the **source code** to the **compiler** (more on this later)
- Think of it as an advanced **search-and-replace**
- For example, the C preprocessor replaces **#include <stdio.h>** with the **entire content** of the header file **stdio.h** (which defines a bunch of utility functions, like **printf**)

More preprocessing

Preprocessor “constants”

- Useful to define **mnemonics** for **numerical constants**. Example:

```
#define TRUE 1  
#define FALSE 0
```

- **Not really constants** (intended as program variables with a fixed value)
- The preprocessor simply **replaces** “TRUE” with “1” and “FALSE” with “0” **throughout the program**

Even more preprocessing

Preprocessor macros

- Using to define **mnemonics** for **short snippets of code**

```
def min(a, b) (a < b ? A : b)
```

- **Incidentally, what's happening in the code above?**
- Useful in the olden days as it is **faster than a function call**
- **Not very useful nowadays** (compilers are very good at **inlining** short functions)

Preprocessor: putting it all together

```
// Preprocessor: example of include
#include <stdio.h>
#include <stdlib.h>

// Preprocessor: example of "constant"
#define TRUE 1
#define FALSE 0

// Preprocessor: example of macro
#define min(a,b) (a < b ? a : b)

int main(int argc, char** argv) {
    int val1 = atoi(argv[1]);
    int val2 = atoi(argv[2]);

    int val3 = min(val1, val2);
    printf("min(%d, %d) is %d\n", val1, val2, val3);

    return 0;
}
```

Back to the parts of a program

```
#include <stdio.h>
```


```
void do_stuff();
```

```
int main() {  
    do_stuff();  
}
```

```
void do_stuff() {  
    printf("Hello, world!");  
}
```

} — ~~Preprocessor directives~~

} — Functions/globals declarations



} — Main function (entry point)

} — Function definitions

Function declaration

Specifies type, name and parameters of a function

- Example #1:

void do_stuff();

The function takes no parameters

Name is "do_stuff"

Return type is "void" (does not return anything)

- Example #2:

int do_other_stuff(float a, float b);

What about this declaration?

Function declarations!

What are they good for?

- **A few things:**
 - C compilers do not allow you to **use a function before it has been defined**
 - You can make sure you always define functions before you call them
 - Can you? What about functions that **call each other**?
- Even if you are not in that case, you may want to use a function somewhere but **define it somewhere else** (e.g., another source code file)
- Function declarations tell the compiler **how the function looks**, so you can use it without having defined it yet

Declaration: a multi-file example

(Do not worry, we'll get to compiling later)

```
C example03_f1.c > main()
1  #include <stdio.h>
2
3  unsigned pow2(unsigned val);
4
5  int main() {
6      unsigned rez = pow2(3);
7      printf("2^3 is %u\n", rez);
8      return 0;
9  }
```

```
C example03_f2.c > pow2(unsigned)
1  unsigned pow2(unsigned val) {
2      int i, rez = 2;
3      if ( val == 0 )
4          return 1;
5      for (i = 1; i < val; i++)
6          rez *= 2;
7      return rez;
8  }
```

```
lorenzo@ensf461:~/class/ensf461F23/lecture04$ gcc -o example03 example03_f1.c example03_f2.c
lorenzo@ensf461:~/class/ensf461F23/lecture04$ ./example03
2^3 is 8
lorenzo@ensf461:~/class/ensf461F23/lecture04$ _
```

Let's take a break and quiz!

Navigate to D2L->Quizzes->Quiz 4

Back to the parts of a program

```
#include <stdio.h>
```

```
void do_stuff();
```

```
int main() {  
    do_stuff();  
}
```

```
void do_stuff() {  
    printf("Hello, world!");  
}
```

} — ~~Preprocessor directives~~

} — ~~Functions/globals declarations~~

} — Main function (entry point) 

} — Function definitions

Function definition: the `main ()` function

What is the difference between declaring and defining?

- A function **declaration** describes the **prototype** of a function
- A function **definition** describes the actual **implementation**
- Each C function must define a function called **`main`** that defines **where execution must start** (and typically end)
- By convention, the function has type **`int`** and its return value “says something” about whether the program **executed correctly** (return **`0`** if everything good 👍, something **`!= 0`** if something went bad/unexpected 🙅)

main () function: examples

```
int main() {  
    return 0;  
}
```

This syntax is used to receive **command line arguments** (let's see an example...)

```
int main(int argc, char** argv) {  
    int i;  
    printf("Number of arguments: %d\n", argc);  
    for (i=0; i<argc; i++)  
        printf("Argument #%d: %s\n", i, argv[i]);  
    return 0;  
}
```

Back to the parts of a program

```
#include <stdio.h>
```

```
void do_stuff();
```

```
int main() {  
    do_stuff();  
}
```

```
void do_stuff() {  
    printf("Hello, world!");  
}
```

} — ~~Preprocessor directives~~

} — ~~Functions/globals declarations~~


} — ~~Main function (entry point)~~

} — Function definitions ←

Function definitions

- Not much more to say
- A function definition is similar to a function declaration, but followed by the actual **function body**
 - Function body is defined within curly braces
 - The return value of the function must match what was defined in its prototype

Function definition example

```
C example03_f2.c >  pow2(unsigned)
1  unsigned pow2(unsigned val) {
2      int i, rez = 2;
3      if ( val == 0 )
4          return 1;
5      for (i = 1; i < val; i++)
6          rez *= 2;
7      return rez;
8  }
```

Let's talk about compilers

C is a **compiled** language

- **What does it mean?**
- Languages like Bash/Python/Node.js are **interpreted**
 - In principle, a program called **interpreter** **reads your code** and **executes** it
 - Under the hood there is a lot of funky stuff going on to make things go faster, but the above abstraction holds in practice
- **C does nothing of that stuff**

What is a compiled language?

- Source code is not executed by an interpreter
- Instead, is passed to a **compiler** which turns it into **executable machine code**
- The output of the compiler is an executable file which can be executed directly

Execution &
return value

The screenshot shows a code editor window titled "example05.c — lecture04 [SSH: 192.168.64.3]". The editor has four tabs: "example03_f1.c", "example03_f2.c", "example04.c", and "example05.c". The "example05.c" tab is active, showing the following source code:

```
1 #include <stdio.h>
2
3 int main() {
4     return 42;
5 }
6
```

A callout box labeled "Source code" points to the source code. Below the source code, the terminal output shows the compilation and execution process:

```
lorenzo@ensf461:~/class/ensf461F23/lecture04$ gcc -o example05 example05.c
lorenzo@ensf461:~/class/ensf461F23/lecture04$ objdump -d example05 | tail -n 12
0000000000000714 <main>:
714: 52800540      mov     w0, #0x2a      // #42
718: d65f03c0      ret

Disassembly of section .fini:
000000000000071c <_fini>:
71c: d503201f      nop
720: a9bf7bfd      stp     x29, x30, [sp, #-16]!
724: 91003fd      mov     x29, sp
728: a8c17bfd      ldp     x29, x30, [sp], #16
72c: d65f03c0      ret

lorenzo@ensf461:~/class/ensf461F23/lecture04$ ./example05 ; echo $?
42
```


A callout box labeled "Compilation" points to the compilation command. Another callout box labeled "Machine code (binary)" points to the disassembly output. A third callout box labeled "Execution & return value" points to the execution command and its output.

How do I do this?

Well, you need a **compiler**

- In this class we are going to use **gcc**
 - The first (and for many years the only) **OSS industry-grade compiler**
 - Ships with pretty much **every Linux distribution**
 - Still the **most popular C compiler**
- You should learn to **run gcc** and **compile programs** from the **command line**

Which compilers do you regularly use?

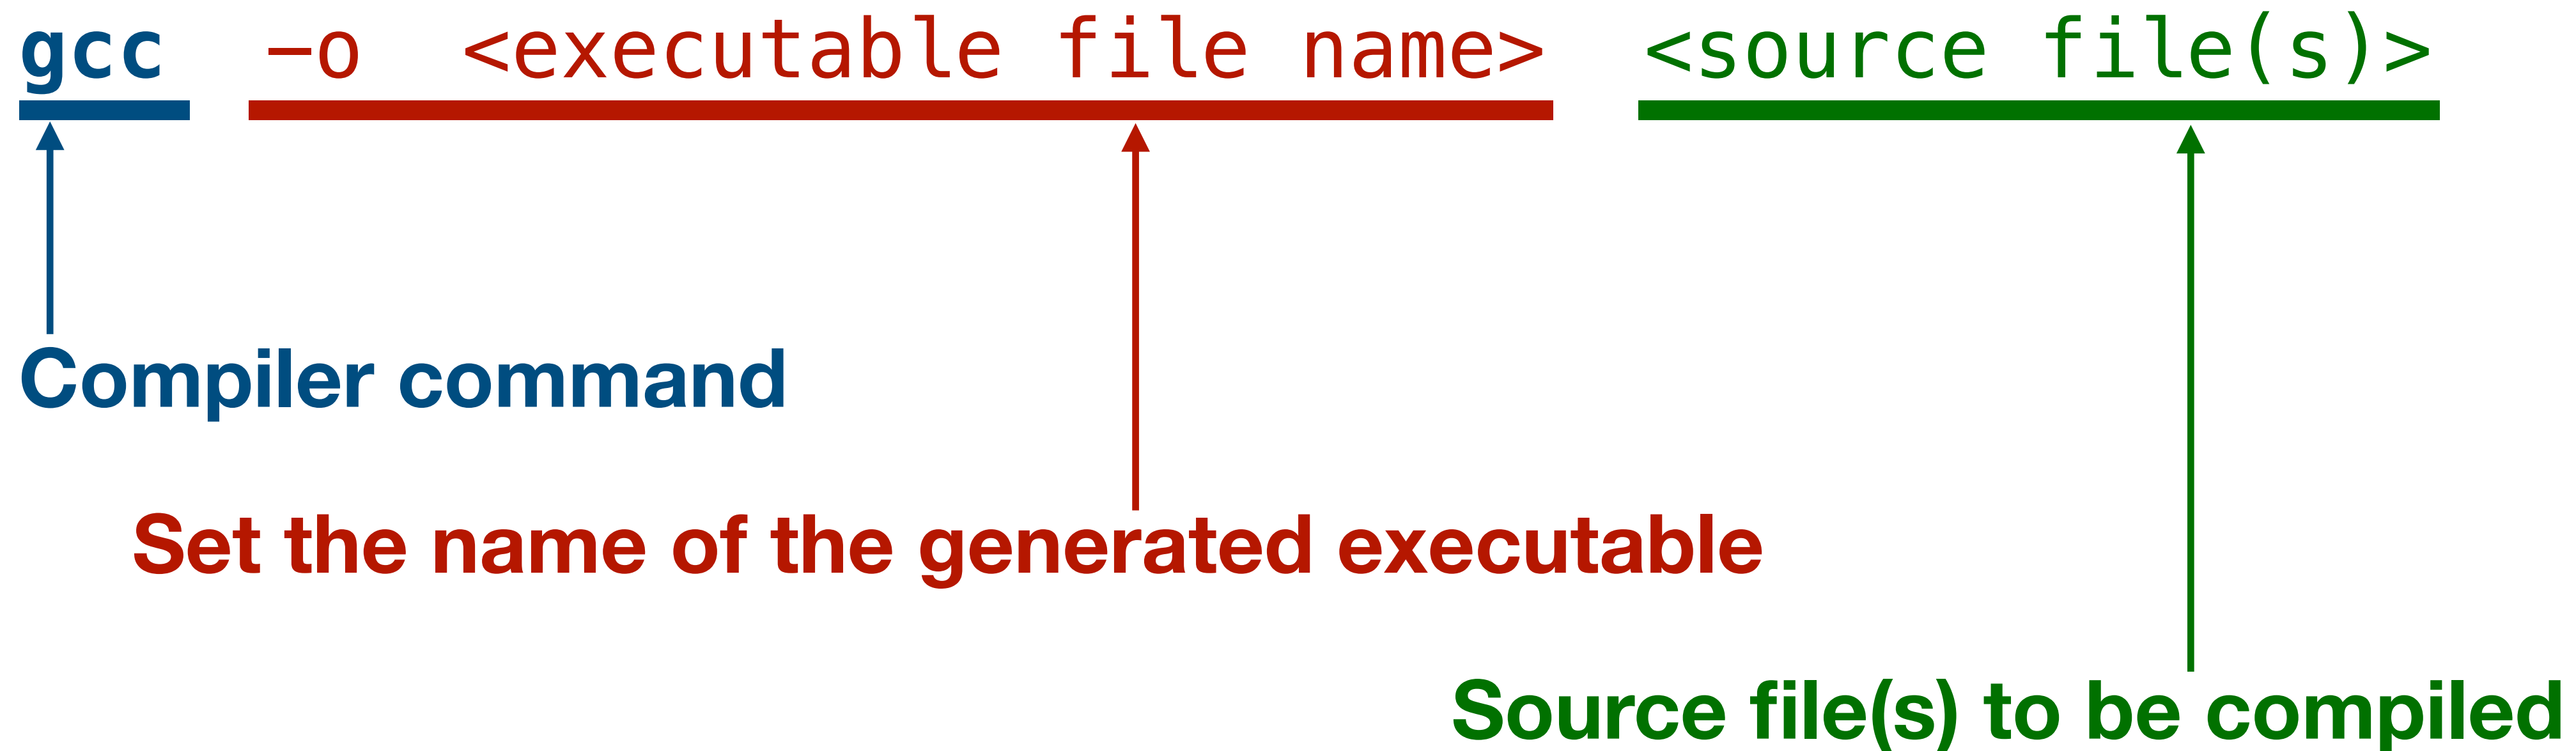


GCC is the most popular compiler, but compilers for microcontrollers in Embedded development are also particularly popular with a share of 38%, which puts them above Clang and MSVC.

<https://www.jetbrains.com/idea/devecosystem-2020/c/>

Running gcc: the basics

In a nutshell, this is how you do it



Source file(s)?

A C program can be split across multiple files

- You can compile such a program with **gcc -o prog file1.c file2.c**
- However, if you define a function in a source file and use it in another one, you need to make the **prototype** of that function known
- You can use **header files** to do this

Header files

```
C example06_f1.c > ...  
1  #include <stdio.h>  
2  #include "example06_f2.h"  
3  
4  int main(int argc, char** argv) {  
5  
6      int val = multiply_by_42(argc);  
7      printf("Value: %d\n", val);  
8  
9      return 0;  
10 }
```

```
C example06_f2.h > ...  
1  int multiply_by_42(int val);  
2
```

```
C example06_f2.c > multiply_by_42(int)  
1  int multiply_by_42(int val) {  
2      return val*42;  
3  }  
4
```

Now for the fun part of C

Pointers & memory management

C memory management & you

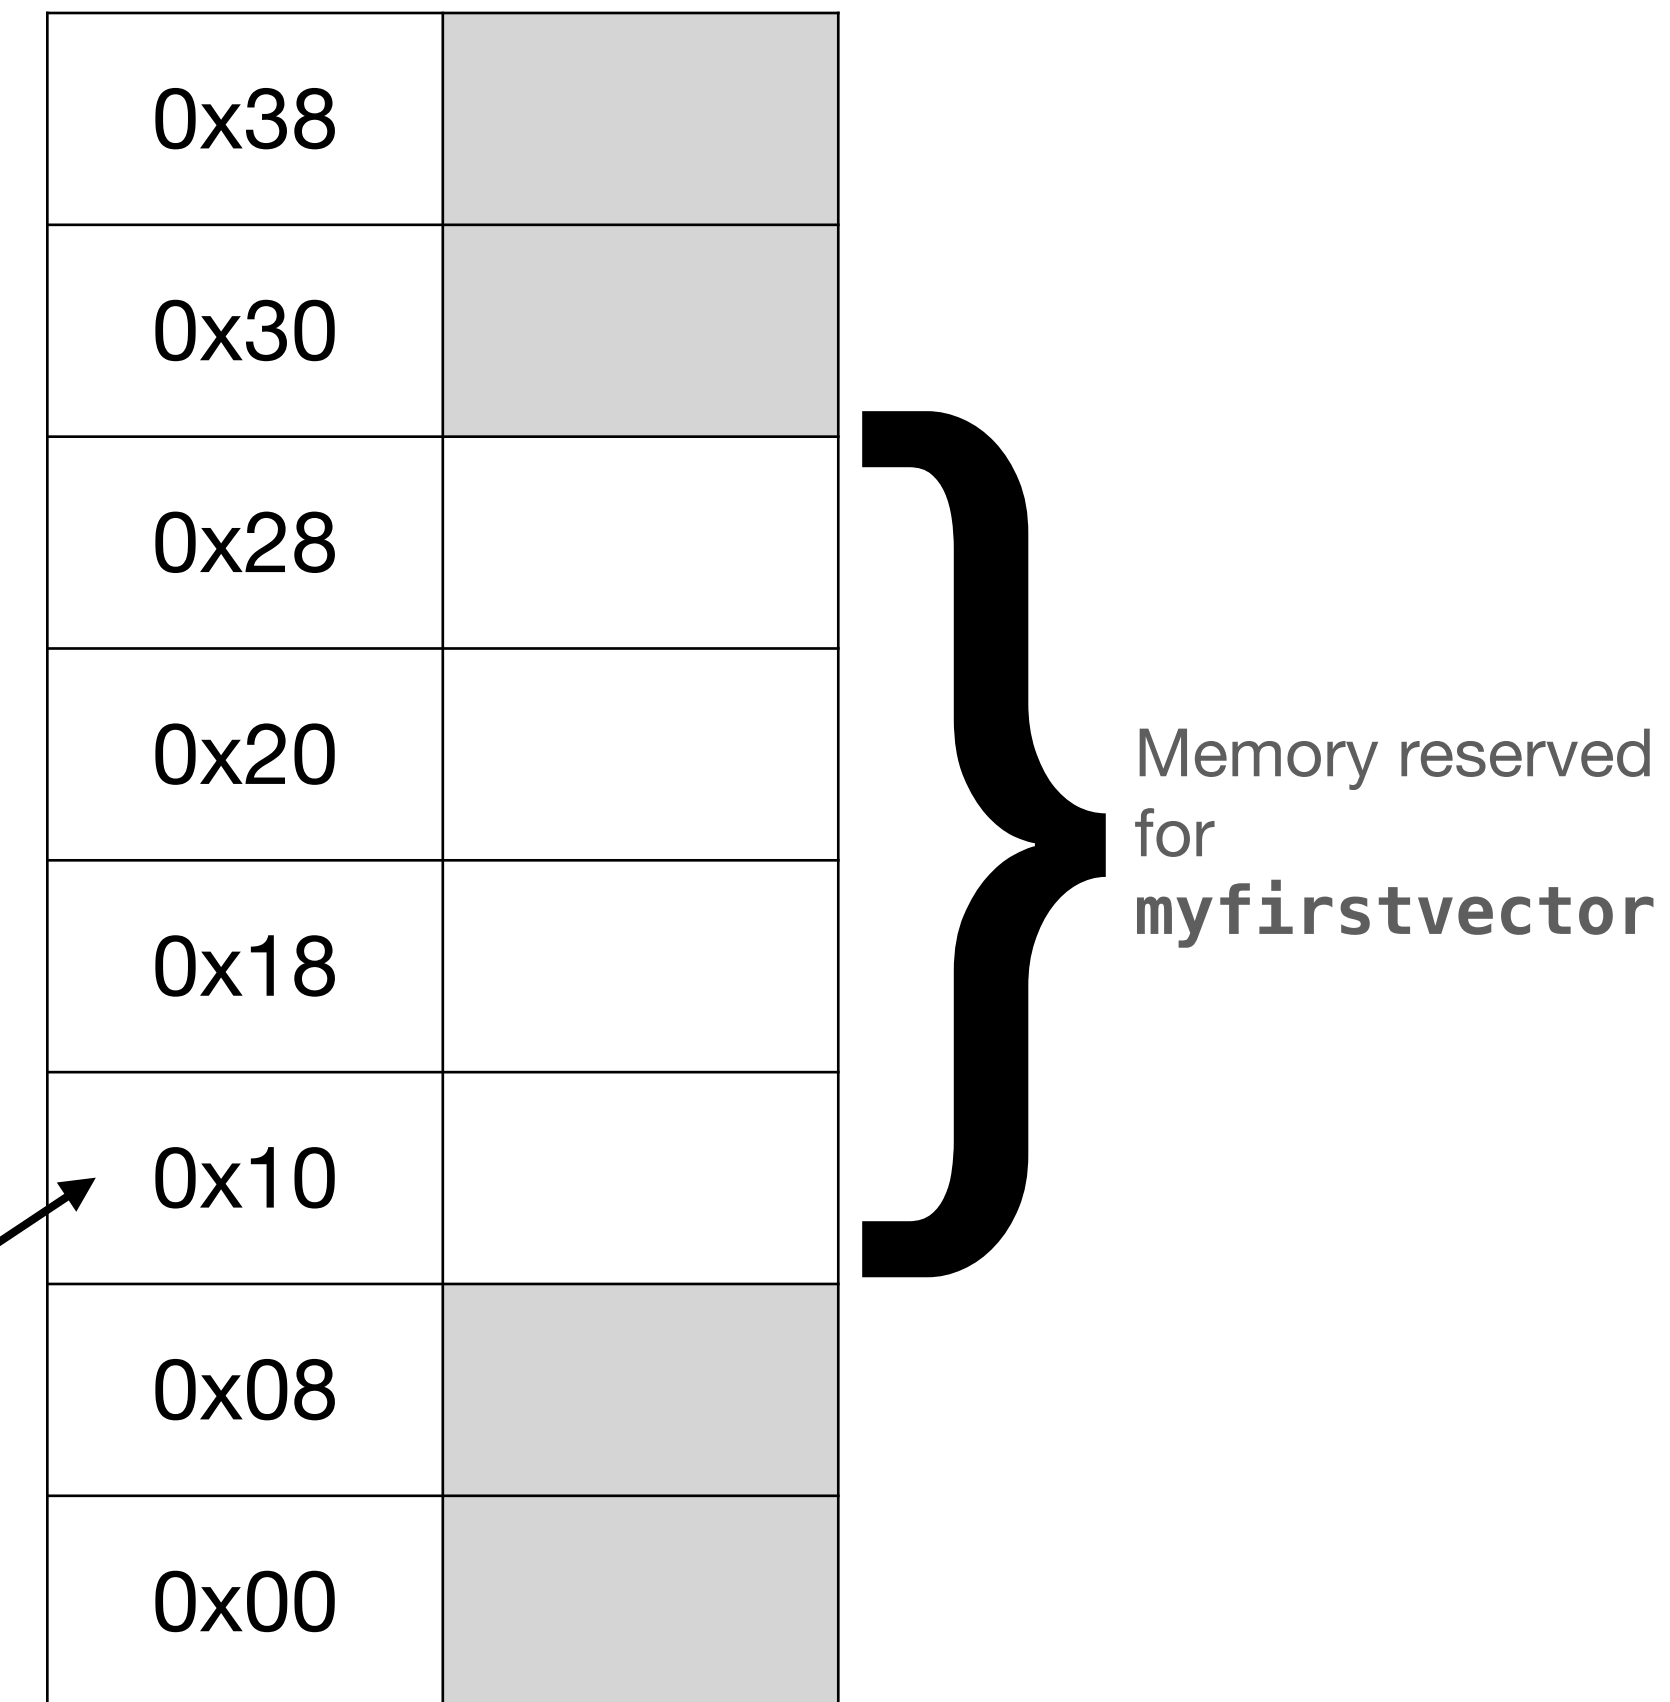
- C, unlike many other languages, allows to **directly access memory locations** where **variables** are stored
- A few caveats:
 - Those are not **physical** memory locations
 - Those are not **absolute** memory locations
 - Those locations may **not remain the same** across executions

Vectors & vector indexing

Let's start with an easy one

- In C a vector is an array of elements of a given type
- Example: `int myfirstvector[4];`
- `myfirstvector` is really a pointer 🤯
 - The location where the array starts in memory

Location pointed by
`myfirstvector`



Vectors/2

- When code accesses an **element** of a vector:

```
myfirstvector[2] = 3;
```

- The compiler computes the following **memory location**:

```
myfirstvector + (3 * sizeof(int))
```

- ...and stores the value “3” in it!

What is this?



Pointers

- A variable storing a **memory location** is a **pointer**

```
int myvar;    // Variable (stores an integer value)
```

```
int *myvar:   // Pointer (stores the value of an  
              // integer-sized memory location)
```


How do I get a memory location anyway?

The “&” operator

- A pointer can store a memory location, but how do you make it point to something useful?
- One way to get the address of a variable is to use the “&” operator:

```
int val = 3;  
int* myptr = &val;
```

- To write at the location pointed at by myptr, use the “*” operator:

```
*myptr = 42;  
printf(“%d\n”, val);
```

[] and * are related

- `int vec[3]` and `int* ptr` are both pointers, but...
- The first guarantees that there is a three-integer memory region allocated starting at the memory location pointed by `var`
- `ptr` is just a pointer! You can store a location in it, but there is no functional guarantee that it is going to point at a useful or reasonable place in memory
- **What happens if you try this?**

```
int* ptr = 0;  
*ptr = 42;
```

Let's look at an example

```
#include <stdio.h>

void value_or_reference(int val, int* ref) {
    val = 42;
    *ref = 42;
}

int main() {
    char myvec[4] = {'a', 'b', 'c', 'd'};

    printf("Position 0 of myvec is %c\n", myvec[0]);
    printf("Position 1 of myvec is %c\n", *(myvec+sizeof(char)));

    int byvalue = 2;
    int byref = 2;
    value_or_reference(byvalue, &byref);

    printf("Variable passed by value: %d\n", byvalue);
    printf("Variable passed by references: %d\n", byref);

    return 0;
}
```

Finally, saving you some time

DON'T DO THIS!

```
gcc -o myprogram.c myprogram.c
```

- **What's wrong with the line above?**
- It passes the **input file** (myprogram.c) as **output file**!

That's all!